



# Model Selection

Azercell DataMinds Bootcamp

DS4

# Content

- Introduction
  - Model selection
- Linear Models
  - Logistic regression
  - Linear regression
- Metrics
  - Classification metrics
    - Acc/Precision/Recall
    - ROC AUC curve
  - Regression Metrics
- Regularization
  - L1 regularization
  - L2 regularization
- Model selection technique
  - Strategy
  - Kfold
  - Hypothesis testing
- Ensemble Methods
  - Bagging
  - Stacking
  - Boosting

# Introduction

# Model Selection Train/Test

Key important features

- Bias/Variance tradeoff
- Efficiency
- Evaluation and comparison

Model Model selection is a **critical step in machine learning** because not all models perform equally well on a given dataset.



Reason	Explanation
Better performance	Different models fit different data well
Generalization	Avoid underfitting and overfitting
Computational cost	Choose efficient models for production
Task-specific needs	Match model to data and problem type
Evaluation rigor	Use metrics and validation to choose reliably

Production pipeline

Data Gathering

EDA

Feature Engineering  
advanced techniques

Model selection: Train/Test

Develop production code

Deploy and monitor

# Linear Models

# Introduction to Linear and Logistic Regression

## Overview

- Both are **supervised learning** algorithms.
- They differ in **output type** and **use case**.

Feature	Linear Regression	Logistic Regression
Target Variable	Continuous	Categorical (binary/multiclass)
Output	Any real number	Probability (0 to 1)
Used For	Regression tasks	Classification tasks

## Common Use Cases:

- Linear Regression: Predicting housing prices, stock trends.
- Logistic Regression: Spam detection, disease prediction.

# Linear Regression – Formula & Intuition

## Model Equation:

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

- $\hat{y}$ : Predicted value
- $\beta_i$ : Coefficients/weights
- $x_i$ : Feature variables

## Loss Function (MSE):


$$MSE = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

**Objective:** Minimize MSE by adjusting weights using **Gradient Descent** or **Closed-form solution**.

## Gradient Descent – Pseudocode.

```
Initialize weights  $\beta_0, \beta_1, \dots, \beta_n$  randomly
Repeat until convergence:
    Predict:  $y_{\text{hat}} = X \cdot \beta$ 
    Compute gradients:  $\nabla = -2/n * X^t \cdot (y - y_{\text{hat}})$ 
    Update weights:  $\beta = \beta - \alpha * \nabla$ 
```

- $\alpha$ : Learning rate
- $X$ : Input matrix
- $y$ : Target vector
- Convergence = when loss stops decreasing significantly.



```
from sklearn.linear_model import LinearRegression
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split
```

```
# Generate sample data
```

```
X, y = make_regression(
    n_samples=100,
    n_features=1,
    noise=10)
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2)
```

```
# Fit linear regression model
```

```
model = LinearRegression()
model.fit(X_train, y_train)
```

```
# Coefficients
```

```
print("Coefficient:", model.coef_)
print("Intercept:", model.intercept_)
```

```
# Predict
```

```
y_pred = model.predict(X_test)
```

## Interpretation:

The coefficient tells how much  $y$  changes for one unit increase in  $X$ . The intercept is the expected  $y$  when all features are 0.

# Logistic Regression – Formula & Intuition

## Model Equation (Sigmoid Function):

$$\hat{p} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

where  $z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$

- $\hat{p}$ : Predicted probability of class 1
- $\sigma(z)$  Sigmoid function maps  $z$  to  $[0,1]$

## Decision Rule:

- If  $p > 0.5$  : predict class 1 (by default)
- Else: predict class 0

## Loss Function (Binary Cross-Entropy):

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{p}_i) + (1 - y_i) \log(1 - \hat{p}_i)]$$

Initialize weights  $\beta_0, \beta_1, \dots, \beta_n$  randomly


Repeat until convergence:

Predict:  $p\_hat = \text{sigmoid}(X \cdot \beta)$

Compute error:  $error = p\_hat - y$

Compute gradients:  $\nabla = (1/n) * X^t \cdot error$

Update weights:  $\beta = \beta - \alpha * \nabla$



```
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
```

```
# Generate binary classification data
```

```
X, y = make_classification(
    n_samples=100,
    n_features=2,
    n_classes=2)
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2)
```

```
# Fit logistic regression model
```

```
model = LogisticRegression()
model.fit(X_train, y_train)
```

```
# Predict
```

```
y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))
```

## Interpretation:

- `model.predict_proba(X_test)` gives class probabilities.
- Useful for threshold tuning, ROC/AUC analysis.

- Gradient is from the derivative of the log-loss with respect to weights.
- Learning rate  $\alpha$  must be tuned to ensure convergence.



# Linear vs Logistic – Summary & Tips

Aspect	Linear Regression	Logistic Regression
Output	Continuous	Probabilistic (0–1)
Loss Function	MSE	Cross-Entropy
Output Function	Identity	Sigmoid
Gradient Formula	$\nabla = -2X^T(y - \hat{y})$	$\nabla = X^T(\sigma(z) - \hat{y})$

## Best Practices:

- Normalize features (important for GD convergence).
- Monitor learning curves during training.
- Use `learning_rate`, `max_iter`, and `tol` to control GD in custom implementations.

# Metrics

# Classification Metrics Overview

## Metrics

**Goal:** Evaluate how well the model distinguishes between classes (binary or multiclass).

### Key Metrics:

- **Precision** – How many selected items are relevant?
- **Recall** – How many relevant items are selected?
- **F1 Score** – Harmonic mean of precision and recall.
- **Confusion Matrix** – Summarizes predictions vs actuals.
- **ROC Curve & AUC** – Probability-based discrimination measure.

```
from sklearn.metrics import precision_score, recall_score, f1_score
```

```
y_true = [0, 1, 1, 1, 0, 1, 0, 0, 1]
y_pred = [0, 1, 0, 1, 0, 1, 1, 0, 1]
```

```
print("Precision:", precision_score(y_true, y_pred))
print("Recall:", recall_score(y_true, y_pred))
print("F1 Score:", f1_score(y_true, y_pred))
```

Precision: 0.8

Recall: 0.8

F1 Score: 0.8

**TP (True Positive)** – Correctly predicted positive.


**FP (False Positive)** – Incorrectly predicted positive.

**TN (True Negative)** – Correctly predicted negative.

**FN (False Negative)** – Incorrectly predicted negative.

```
from sklearn.metrics import (confusion_matrix,
                             ConfusionMatrixDisplay)
import matplotlib.pyplot as plt
```

```
cm = confusion_matrix(y_true, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot()
plt.show()
```



# Classification Metrics definition

## Metrics

Metric	Formula	Meaning
<b>Accuracy</b>	$\frac{TP + TN}{TP + TN + FP + FN}$	Overall correctness
<b>Precision</b>	$\frac{TP}{TP + FP}$	Correctly predicted positives
<b>Recall</b>	$\frac{TP}{TP + FN}$	Found all actual positives

**High Accuracy** is misleading on imbalanced data.

**Precision** is crucial when **FP is costly** (e.g., spam filters).

**Recall** is vital when **FN is dangerous** (e.g., cancer screening).

**TP (True Positive)** – Correctly predicted positive.

**FP (False Positive)** – Incorrectly predicted positive.

**TN (True Negative)** – Correctly predicted negative.

**FN (False Negative)** – Incorrectly predicted negative.




# ROC Curve & AUC — Threshold-Free Performance

## Metrics

- **TPR (Recall):**  $\frac{TP}{TP+F}$
- **FPR:**  $\frac{FP}{FP+TN}$
- **ROC Curve:** TPR vs FPR at different thresholds
- **AUC (Area Under Curve):** Scalar summary of classifier's discrimination ability
  - **AUC = 0.5** → Random guessing
  - **AUC = 1.0** → Perfect classifier

## Why Use ROC AUC?

- Unaffected by class imbalance
- Evaluates performance **across all thresholds**



```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, roc_auc_score
```

```
# Logistic Regression
model = LogisticRegression()
model.fit(X_train, y_train)
# Predict probabilities
y_scores = model.predict_proba(X_test)[:, 1]
```

```
# ROC and AUC
fpr, tpr, _ = roc_curve(y_test, y_scores)
auc_score = roc_auc_score(y_test, y_scores)
```

```
plt.figure(figsize=(6, 4))
plt.plot(fpr, tpr, label=f"AUC = {auc_score:.2f}")
plt.plot([0, 1], [0, 1], linestyle="--", color="gray")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve – Titanic Dataset")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

# Regression Metrics

## Metrics

Metric	Formula	Description & Use Case
<b>MAE</b>	$MAE = \frac{1}{n} \sum$	$y_i - \hat{y}$
<b>MSE</b>	$MSE = \frac{1}{n} \sum_{i=1}^n (y_t - \hat{y}_i)^2$	Penalizes large errors more than MAE. Use when large deviations are critical.
<b>RMSE</b>	$RMSE = \sqrt{MSE}$	Same units as the target variable. Useful for interpretability.
<b>R<sup>2</sup> Score</b>	$R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$	Proportion of variance in target explained by model. Use to assess goodness of fit. $R^2 = 1$ : perfect fit, $R = 0$ , $R^2 = 0$ : no explanatory power.

**MAE:** Simple interpretation, less sensitive to outliers

**MSE:** Heavier penalty on large errors

**RMSE:** Useful for expressing error in original units

**R<sup>2</sup>:** Good for comparing model fits, but can be misleading with non-linear data

```
from sklearn.metrics import (  
    mean_absolute_error,  
    mean_squared_error,  
    r2_score)
```



# Regularization

# Why do we need Regularization?

## Regularization

### Problem: Overfitting in Machine Learning


- **Overfitting:** The model performs well on training data but poorly on unseen data.
- Happens when:
  - Too complex model (e.g., high-degree polynomial)
  - Too many parameters or not enough data

### Solution: Regularization

- **Idea:** Penalize large model parameters to reduce complexity and prevent overfitting.

$$\text{Cost Function (Basic)} = \text{Loss}(y, \hat{y}) = \sum (y_i - \hat{y}_i)^2$$

- Regularization adds a **penalty** term to this cost function.



## Types of Regularization (L1 & L2)

### L2 Regularization (Ridge Regression)

$$J(\theta) = \sum (y_i - \hat{y}_i)^2 + \lambda \sum (\theta_j^2)$$

- Penalizes **squared** weights
- Shrinks weights smoothly toward zero

### L1 Regularization (Lasso Regression)

$$J(\theta) = \sum (y_i - \hat{y}_i)^2 + \lambda \sum |\theta_j|$$

Can shrink some weights **exactly to zero** → feature selection

**$\lambda$  (lambda)** controls penalty strength:

- $\lambda = 0 \rightarrow$  no regularization
- $\lambda \rightarrow \infty \rightarrow$  all weights shrink heavily



# Regularization example

## Regularization

```
from sklearn.linear_model import Ridge, Lasso
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split
```

```
X, y = make_regression(n_samples=100, n_features=20, noise=15)
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

# Ridge (L2)


```
ridge = Ridge(alpha=1.0)
ridge.fit(X_train, y_train)
print("Ridge R^2:", ridge.score(X_test, y_test))
```

# Lasso (L1)

```
lasso = Lasso(alpha=0.1)
lasso.fit(X_train, y_train)
print("Lasso R^2:", lasso.score(X_test, y_test))
```

Ridge R<sup>2</sup>: 0.9873922504549839

Lasso R<sup>2</sup>: 0.9877459836451575



# Model selection

# Starting with a Baseline Model

Establishing the Baseline in Model Selection

## Purpose of Baseline:

- A baseline model serves as a reference point to compare the performance of more complex models.
- It is simple, fast to train, and interpretable.

## Examples of Baseline Models:


- **DummyClassifier** (e.g., always predicts the most frequent class).
- **Linear Regression** without feature scaling or regularization.
- **Mean or median predictor** for regression.

## Why Start with a Baseline?

- **Sanity check:** Ensures your data pipeline is working correctly.
- **Benchmark:** Helps understand whether complex models actually add value.
- **Interpretability:** Provides initial insights into the task difficulty.

A good model should outperform the baseline significantly. If it doesn't, either the model is underfitting, or the data is not informative enough.

Model Comparison



# Classical ML Models and Performance Estimation

## Exploring Classical Machine Learning Models

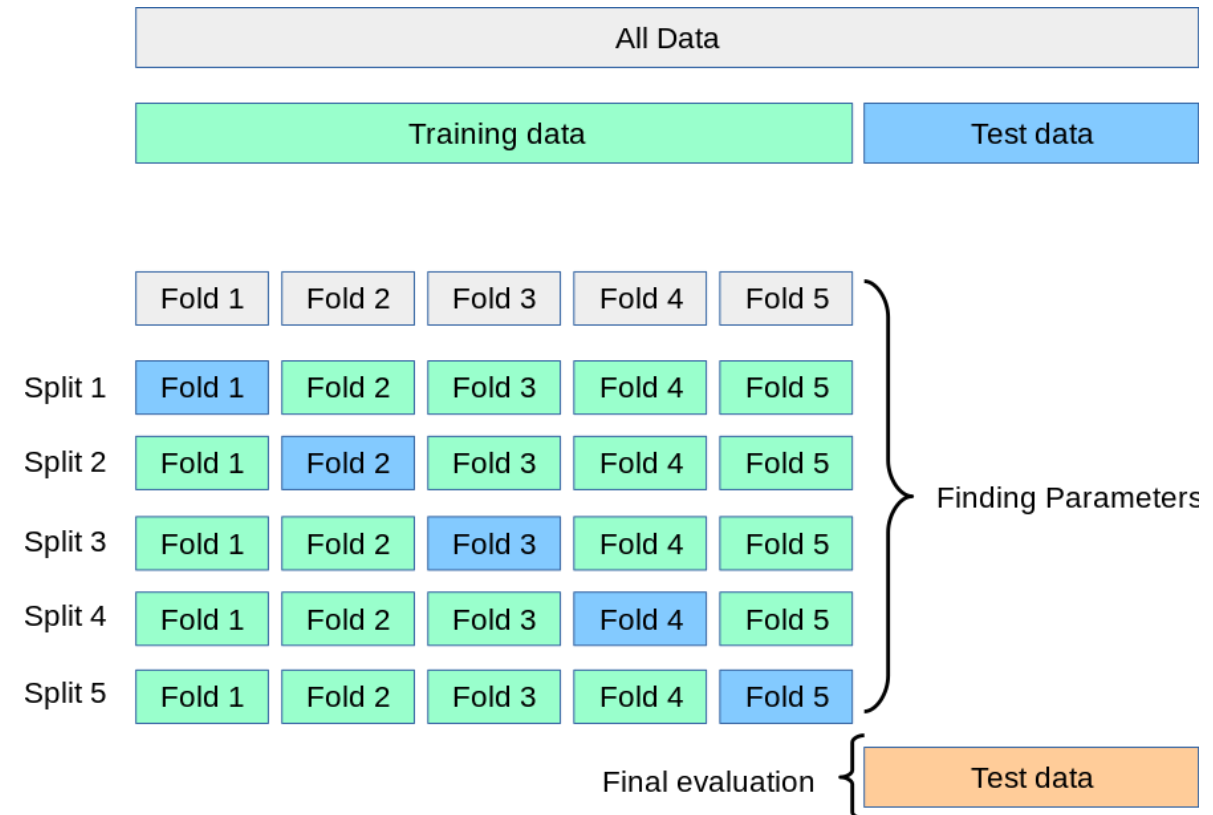
### Candidate Models:

- **Logistic Regression** – linear decision boundaries, interpretable, good for linearly separable data.
- **K-Nearest Neighbors (KNN)** – non-parametric, sensitive to feature scaling.
- **Decision Trees** – interpretable, but prone to overfitting.
- **Random Forest** – ensemble of trees, reduces overfitting, less interpretable.
- **Support Vector Machines (SVM)** – effective in high-dimensional spaces, uses kernel tricks.

### Model Training and Evaluation:

- **Cross-validation (e.g., k-fold)** ensures robust performance estimation.
- Evaluate with metrics:
  - Classification: accuracy, precision, recall, F1-score, AUC-ROC.
  - Regression: MSE, RMSE, MAE,  $R^2$ .

*Bias-variance tradeoff guides model choice. Simpler models may underfit (high bias), complex ones may overfit (high variance). Regularization can help balance this.*



# Iterative Model Refinement and Selection

Refining, Tuning, and Selecting the Optimal Model

## Model Tuning:

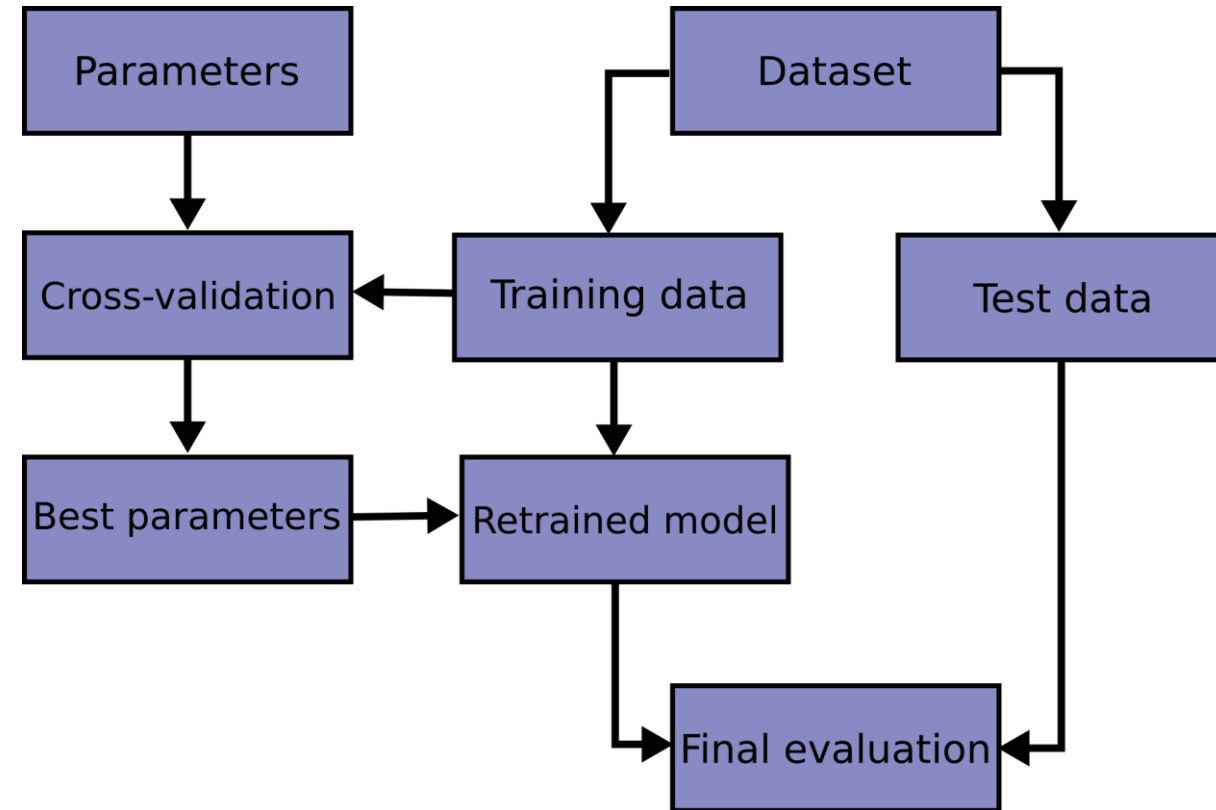
- **GridSearchCV / RandomizedSearchCV** to optimize hyperparameters.
- Use validation performance to avoid overfitting to training data.
- Apply **feature scaling, encoding, or dimensionality reduction** as needed.

## Model Comparison:

- Compare performance using **statistical tests** (e.g., paired t-test) over cross-validation folds.
- Consider **interpretability vs. performance** tradeoff.

## Final Selection Criteria:

- Best generalization performance on **unseen test data**.
- Meets application-specific constraints (e.g., inference speed, memory usage).
- Model stability and explainability.



*•The no free lunch theorem tells us that no model is best for all tasks — model selection is inherently data-dependent.*

# Model selection example

## Model selection

```
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from scipy.stats import ttest_rel
import numpy as np
```

```
# Load data
X, y = load_breast_cancer(return_X_y=True)
```

```
# Define models
model1 = LogisticRegression(max_iter=1000, solver='liblinear') # Simpler linear model
model2 = RandomForestClassifier(random_state=42) # More complex ensemble
```

```
# Stratified K-Folds
cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=1)
```

```
# Evaluate both models
scores1 = cross_val_score(model1, X, y, cv=cv, scoring='accuracy')
scores2 = cross_val_score(model2, X, y, cv=cv, scoring='accuracy')
```

```
print("Logistic Regression Accuracies:", scores1)
print("Random Forest Accuracies:", scores2)
```

```
Logistic Regression Accuracies: [0.94736842
0.94736842 0.92982456 0.94736842 0.92982456
0.96491228 0.98245614 0.96491228 0.96491228
0.96428571]
```

```
Random Forest Accuracies: [0.94736842
0.94736842 1. 0.98245614 0.9122807
0.94736842 1. 0.94736842 0.94736842
0.92857143]
```

- **Paired t-test** offers a statistically principled way to compare models.
- **Cross-validation** ensures fair and robust performance estimation.
- Use **scientific significance (p-value)** alongside **practical significance (accuracy improvement)** when choosing the best model.

```
# Paired t-test
t_stat, p_value = ttest_rel(scores2, scores1)
```

```
print(f"t-statistic: {t_stat:.4f}")
print(f"p-value: {p_value:.4f}")
```

```
t-statistic: 0.1697
p-value: 0.8690
```

- **p-value < 0.05** → We **reject the null hypothesis**: there's a **statistically significant difference** between models.
- **t-statistic > 0** → RandomForestClassifier has **significantly higher accuracy** than LogisticRegression.

## Why Random Forest Might Perform Better

Aspect	Logistic Regression	Random Forest
Model type	Linear classifier	Non-linear ensemble
Assumptions	Assumes linear decision boundaries	Captures complex interactions
Interpretability	High	Moderate to low
Overfitting risk	Low	Moderate (but mitigated via bagging)
Feature importance	Coefficients	Tree-based impurity scores

# Model selection Full pipeline

## Model selection

```
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import KFold, cross_validate
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.compose import ColumnTransformer
import pandas as pd
import numpy as np
```

### # Load dataset

```
X_raw, y = load_breast_cancer(return_X_y=True)
X = pd.DataFrame(X_raw, columns=load_breast_cancer().feature_names)
```

### # Introduce missing values to simulate real-world data

```
X.iloc[:,50, 0] = np.nan
```

### # Preprocessing pipeline for numeric features

```
numeric_features = X.columns.tolist()
numeric_transformer = Pipeline([
    ('imputer', SimpleImputer(strategy='mean')),
    ('scaler', StandardScaler())
])
```

### # ColumnTransformer to apply to all features

```
preprocessor = ColumnTransformer([
    ('num', numeric_transformer, numeric_features)
])
```

### # Full pipeline: preprocessing + model

```
pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', LogisticRegression(max_iter=1000))
])
```

### # Scoring metrics

```
scoring = [
    'accuracy',
    'precision',
    'recall',
    'f1',
    'roc_auc']
```

### # KFold definition

```
kf = KFold(n_splits=5, shuffle=True, random_state=42)
```

### # Evaluate multiple metrics

```
cv_results = cross_validate(
    pipeline,
    X,
    y,
    cv=kf,
    scoring=scoring,
    return_train_score=False)
```

### # Print results

```
for metric in scoring:
    test_scores = cv_results[f'test_{metric}']
    print(f"{metric.capitalize()} scores: {test_scores}")
    print(f"Mean {metric}: {test_scores.mean():.4f}\n")
```

Accuracy scores: [0.97368421 0.98245614 0.96491228 0.99122807 0.97345133] Mean accuracy: 0.9771

Precision scores: [0.97222222 1. 0.95890411 0.98611111 0.97058824] Mean precision: 0.9776

Recall scores: [0.98591549 0.97402597 0.98591549 1. 0.98507463] Mean recall: 0.9862

F1 scores: [0.97902098 0.98684211 0.97222222 0.99300699 0.97777778] Mean f1: 0.9818

Roc\_auc scores: [0.99737963 1. 0.98427776 0.99770717 0.99415964] Mean roc\_auc: 0.9947

# Ensemble methods



# Introduction to Ensemble Learning

## Ensemble methods

### What is Ensemble Learning?

- Ensemble learning combines multiple models (often called **weak learners**) to build a more robust and accurate model.
- It leverages the **wisdom of the crowd**: aggregating multiple predictions typically outperforms a single model.

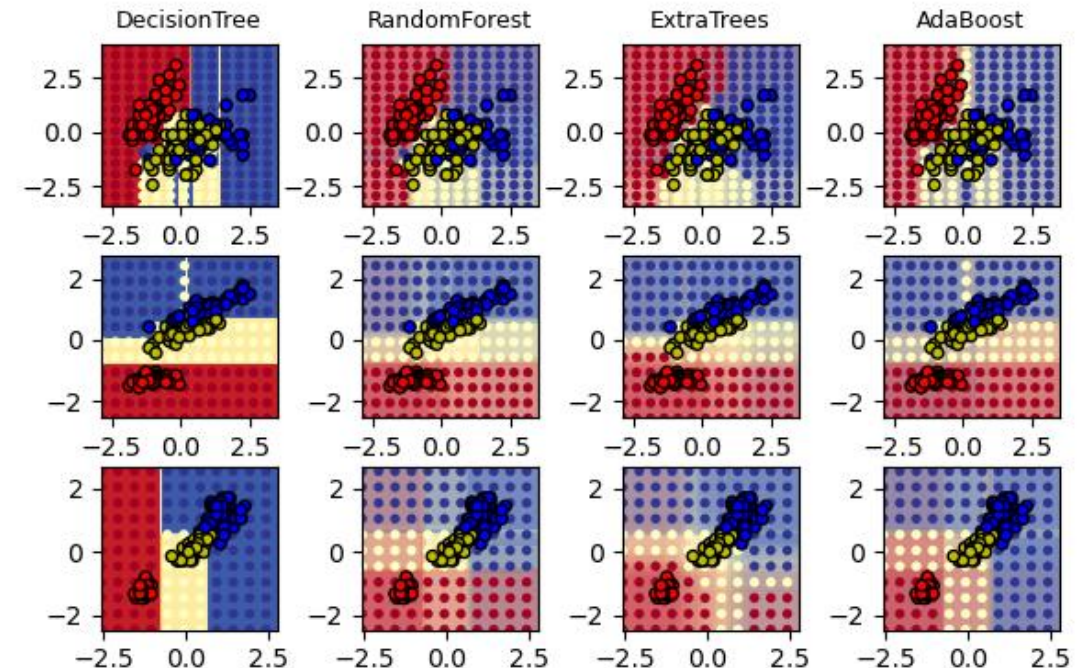
### Types of Ensembles:

1. *Bagging (Bootstrap Aggregating)*
2. *Boosting*
3. *Stacking*

### Why Use Ensembles?

- Reduce variance (Bagging)
- Reduce bias (Boosting)
- Leverage model diversity (Stacking)

Classifiers on feature subsets of the Iris dataset



# Bagging (Bootstrap Aggregating)

Ensemble methods

## Concept:

- Train multiple models on **bootstrapped samples** of data.
- Combine predictions using **voting (classification)** or **averaging (regression)**.

## Pros:

- Reduces variance
- Handles overfitting well

## Cons:

- Limited bias reduction

**Popular Algorithm:** RandomForestClassifier

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
```

```
X, y = load_iris(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
model = RandomForestClassifier(
    n_estimators=100,
    random_state=42)
model.fit(X_train, y_train)
print(model.score(X_test, y_test))
```

# Bagging pipeline

## Ensemble methods

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
```

```
accuracy = make_scorer(accuracy_score)
```

```
cv = KFold(5, shuffle=True, random_state=0)
```

```
model = BaggingClassifier(
    estimator=DecisionTreeClassifier(),
    n_estimators=300,
    bootstrap=True,
    max_samples=1.0,
    max_features=1.0,
    random_state=0)
```

```
column_transform = ColumnTransformer([
    ('categories', categorical_onehot_encoding, low_card_categorical),
    ('numeric', numeric_passthrough, continuous)
],
remainder='drop',
verbose_feature_names_out=False,
sparse_threshold=0.0)
```

```
model_pipeline = Pipeline([
    ('processing', column_transform),
    ('modeling', model)
])
```

```
cv_scores = cross_validate(
    estimator=model_pipeline,
    X=data,
    y=target_median,
    scoring=accuracy,
    cv=cv,
    return_train_score=True,
    return_estimator=True
)
```

```
mean_cv = np.mean(cv_scores['test_score'])
std_cv = np.std(cv_scores['test_score'])
fit_time = np.mean(cv_scores['fit_time'])
score_time = np.mean(cv_scores['score_time'])
```

```
print(f"mean_cv: {mean_cv:0.3f} (std_cv: {std_cv:0.3f})",
      f"fit: {fit_time:0.2f}",
      f"secs pred: {score_time:0.2f} secs")
```

# Boosting

## Ensemble methods

### Concept:

- Train models **sequentially**. Each new model focuses on **correcting the errors** of the previous one.
- Models are **weighted** in the final prediction.

### Pros:


- Reduces both bias and variance
- High accuracy

### Cons:

- Can overfit if not properly regularized
- Training is sequential (slower)

### Popular Algorithms:

- AdaBoost,
- GradientBoosting,
- XGBoost,
- LightGBM
- Catboost



```
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
```

```
# Updated keyword: use 'estimator' instead of 'base_estimator'
model = AdaBoostClassifier(
    estimator=DecisionTreeClassifier(max_depth=1),
    n_estimators=50,
    learning_rate=1.0,
    random_state=42
)
```

```
model.fit(X_train, y_train)
print(model.score(X_test, y_test))
```

0.9473684210526315

# Boosting algorithm

## Ensemble methods

```
from sklearn.tree import DecisionTreeRegressor
import numpy as np
```

```
class GradientBoosting():
    def __init__(self, learning_rate=0.1, n_estimators=10, **params):
        self.learning_rate = learning_rate
        self.n_estimators = n_estimators
        self.params = params
        self.trees = list()

    def sigmoid(self, x):
        x = np.clip(x, -100, 100)
        return 1 / (1 + np.exp(-x))

    def logit(self, x, eps=1e-6):
        xp = np.clip(x, eps, 1-eps)
        return np.log(xp / (1 - xp))

    def gradient(self, y_true, y_pred):
        gradient = y_pred - y_true
        return gradient

    def fit(self, X, y):
        self.init = self.logit(np.mean(y))
        y_pred = self.init * np.ones((X.shape[0],))
        for k in range(self.n_estimators):
            gradient = self.gradient(self.logit(y), y_pred)
            tree = DecisionTreeRegressor(**self.params)
            tree.fit(X, -gradient)
            self.trees.append(tree)
            y_pred += (self.learning_rate * tree.predict(X))
```

```
def predict_proba(self, X):
    y_pred = self.init * np.ones((X.shape[0],))
    for tree in self.trees:
        y_pred += (self.learning_rate * tree.predict(X))
    return self.sigmoid(y_pred)

def predict(self, X, threshold=0.5):
    proba = self.predict_proba(X)
    return np.where(proba >= threshold, 1, 0)
```

# Stacking

## Ensemble methods

### Concept:

- Combine predictions of **multiple base models** using a **meta-model** (often logistic regression or linear regression).
- Base models learn the data, the meta-model learns from their predictions.

### Pros:

- Can combine diverse models (tree + SVM + neural net)
- Captures both low and high bias/variance patterns

### Cons:

- More complex, prone to overfitting if improperly validated

---

#### Algorithm 19.7 Stacking

---

**Input:** Training data  $\mathcal{D} = \{\mathbf{x}_i, y_i\}_{i=1}^m$  ( $\mathbf{x}_i \in \mathbb{R}^n, y_i \in \mathcal{Y}$ )

**Output:** An ensemble classifier  $H$

- 1: Step 1: Learn first-level classifiers
  - 2: **for**  $t \leftarrow 1$  to  $T$  **do**
  - 3:     Learn a base classifier  $h_t$  based on  $\mathcal{D}$
  - 4: **end for**
  - 5: Step 2: Construct new data sets from  $\mathcal{D}$
  - 6: **for**  $i \leftarrow 1$  to  $m$  **do**
  - 7:     Construct a new data set that contains  $\{\mathbf{x}'_i, y_i\}$ , where  $\mathbf{x}'_i = \{h_1(\mathbf{x}_i), h_2(\mathbf{x}_i), \dots, h_T(\mathbf{x}_i)\}$
  - 8: **end for**
  - 9: Step 3: Learn a second-level classifier
  - 10: Learn a new classifier  $h'$  based on the newly constructed data set
  - 11: **return**  $H(\mathbf{x}) = h'(h_1(\mathbf{x}), h_2(\mathbf{x}), \dots, h_T(\mathbf{x}))$
- 

```
from sklearn.ensemble import StackingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
```

```
estimators = [
    ('svc', SVC(probability=True)),
    ('nb', GaussianNB())
]
```

```
model = StackingClassifier(
    estimators=estimators,
    final_estimator=LogisticRegression()
)
model.fit(X_train, y_train)
print(model.score(X_test, y_test))
```

# When to Use Which?


Ensemble methods

Strategy	Use When...	Key Focus
Bagging	High variance models (e.g., decision trees)	Variance
Boosting	High bias models, need higher accuracy	Bias + Errors
Stacking	Need to combine diverse models	Diversity

Bagging → Parallel training (good for speed)

Boosting → Sensitive to outliers (careful with noise)

Stacking → Needs good validation to avoid overfitting



# Hyperparameters tuning



# Grid Search – Exhaustive Hyperparameter Tuning

Hyperparameters tuning

## Definition:

Grid Search tests **all possible combinations** of predefined hyperparameter values.

## How it works:

- You define a **parameter grid** (e.g., {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf']}).
- For each combination, the model is trained and evaluated using **cross-validation**.
- The best combination is selected based on a **scoring metric** (e.g., accuracy, F1).

## Pros:

- Simple and **deterministic**.
- Good for **small search spaces**.

## Cons:

- **Computationally expensive** for large grids.
- Doesn't prioritize promising areas of search space.

```
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC
```

```
param_grid = {
    'C': [0.1, 1, 10],
    'kernel': ['linear', 'rbf']
}
```

```
grid = GridSearchCV(
    SVC(),
    param_grid,
    scoring='accuracy',
    cv=5)
```

```
grid.fit(X_train, y_train)
print(grid.best_params_)
```

```
{'C': 1, 'kernel': 'linear'}
```

# Random Search – Efficient Exploration

Hyperparameters tuning

## Definition:

Random Search randomly samples hyperparameter combinations from given distributions.

## How it works:

- You define **ranges or distributions** (e.g., uniform, log-uniform).
- Random combinations are sampled for a fixed number of iterations.

## Pros:

- **More efficient** than Grid Search when some hyperparameters are more important.
- Useful when there are many parameters.

## Cons:

- Still **requires many evaluations**.
- No feedback from previous trials (i.e., no learning from past evaluations).

```
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import uniform
from sklearn.ensemble import RandomForestClassifier
from scipy.stats import randint
```

```
param_dist = {
    'n_estimators': [50, 100, 200],
    'max_depth': randint(1, 10)
}
rand_search = RandomizedSearchCV(
    RandomForestClassifier(),
    param_dist,
    n_iter=10,
    scoring='accuracy',
    cv=5)
rand_search.fit(X_train, y_train)
print(rand_search.best_params_)
```

```
{'max_depth': 9,
 'n_estimators': 200}
```

# Optuna – Intelligent Optimization

Hyperparameters tuning

Optuna is a **state-of-the-art automatic hyperparameter optimization** framework using **Bayesian optimization and pruning**.

**Key concepts:**

- Uses **TPE (Tree-structured Parzen Estimator)** for smarter search.
- Learns from past trials to **focus on promising regions**.
- Supports **early stopping (pruning)** to skip unpromising trials.

**Basic flow:**

1. Define an `objective(trial)` function.
2. Use Optuna to minimize or maximize the objective.

```
pip install optuna
```



```
import optuna
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score

def objective(trial):
    n_estimators = trial.suggest_int('n_estimators', 50, 200)
    max_depth = trial.suggest_int('max_depth', 2, 32)
    clf = RandomForestClassifier(
        n_estimators=n_estimators,
        max_depth=max_depth)
    return -cross_val_score(
        clf,
        X_train,
        y_train,
        cv=3,
        scoring='accuracy').mean()

study = optuna.create_study(direction='minimize')
study.optimize(objective, n_trials=50)
print(study.best_params)
```

**Why do we use negative cross\_val\_score?**

- `cross_val_score()` returns **positive values** for metrics like accuracy.
- But Optuna requires you to **minimize** or **maximize** your objective explicitly.
- If you want to **maximize accuracy**, you either:
  - **maximize** the score directly, or
  - **minimize the negative** of the score.

# Optuna advanced features

## Hyperparameters tuning

### Advanced capabilities:



- **Pruning:** Stops poor trials early to save time.
- **Visualization:** Provides importance plots, optimization history, and parameter correlation plots.
- **Parallel execution:** Easily run multiple trials in parallel.
- **Sampler customization:** Use TPE, CMA-ES, or random sampling.




```
!pip install nbformat
```

### Visualization example:


```
optuna.visualization.plot_optimization_history(study).show()  
optuna.visualization.plot_param_importances(study).show()
```




# Literature and links



**Ensemble Methods for Machine Learning**  
by [Gautam Kunapuli](#)



**Hands-On Machine Learning with Scikit-Learn**  
by [Aurélien Géron](#)



**Machine Learning for Tabular Data: XGBoost, Deep Learning, and AI**  
by [Mark Ryan Luca Massaron](#)

Kfold

[https://scikit-learn.org/stable/modules/cross\\_validation.html](https://scikit-learn.org/stable/modules/cross_validation.html)

Hyperparameters tune

<https://optuna.org/>

Ensemble methods

stacking:

[https://medium.com/@brijesh\\_soni/stacking-to-improve-model-performance-a-comprehensive-guide-on-ensemble-learning-in-python-9ed53c93ce28](https://medium.com/@brijesh_soni/stacking-to-improve-model-performance-a-comprehensive-guide-on-ensemble-learning-in-python-9ed53c93ce28)

[https://rasbt.github.io/mlxtend/user\\_guide/classifier/StackingClassifier/](https://rasbt.github.io/mlxtend/user_guide/classifier/StackingClassifier/).

bagging

<https://www.datacamp.com/tutorial/what-bagging-in-machine-learning-a-guide-with-examples>



**Thank you for your attention**